

Background

- Processes can execute concurrently
 - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Illustration of the problem:
 - Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
 - We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

in
out

empty
0

full
N

Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

Race Condition

- **counter++** could be implemented as

```
1 register1 = counter
2 register1 = register1 + 1
3 counter = register1
```

- **counter--** could be implemented as

```
4 register2 = counter
5 register2 = register2 - 1
6 counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute **register1 = counter**
 S1: producer execute **register1 = register1 + 1**
 S2: consumer execute **register2 = counter**
 S3: consumer execute **register2 = register2 - 1**
 S4: producer execute **counter = register1**
 S5: consumer execute **counter = register2**

{register1 = 5}
 {register1 = 6}
 {register2 = 5}
 {register2 = 4}
 {counter = 6}
 {counter = 4}

5 ++ -- 5

```

Counter ++
mov 0x8049alc, %eax
add $0x1, %eax
mov %eax, 0x8049alc

```

Counter = 50

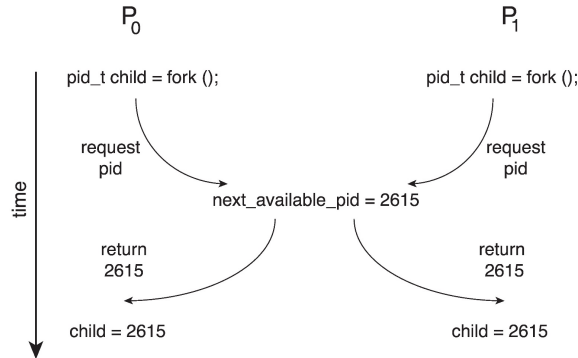
OS	Thread 1	Thread 2	%eax	Counter
			0	50
	mov 0x8049alc, %eax		50	50
	add \$0x1, %eax		51	50
Interrupt	save T1's state			
	restore T2's state		0	50
		mov 0x8049alc, %eax	50	50
		add \$0x1, %eax	51	50
		mov %eax, 0x8049alc	51	51
Interrupt	save T2's state			
	restore T1's state		51	51
	mov %eax, 0x8049alc		51	51

Race Condition (cont.)

- Where multiple processes are writing or reading some shared data and the final result depends on who runs precisely when, are called **race conditions**

Race Condition (cont.)

- Processes P_0 and P_1 are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)




- Unless there is mutual exclusion, the same pid could be assigned to two different processes!

Critical Section & Mutual Exclusion

- Critical Section
 - A section of code in which the process accesses and modifies ***shared variables***
- Mutual Exclusion
 - A method for ensuring that one (or a specified number) of processes are in a critical section

n

Mutual Exclusion



Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc.
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

Critical Section

- General structure of process P_i

```
do {  
    → entry section  
        critical section  
    → exit section  
        remainder section  
} while (true);
```

Algorithm for Process P_i

```
do {  
  entry → while (turn == i);  
           critical section  
  exit  → turn = j;  
           remainder section  
} while (true);
```

Handwritten annotations:

- A red circle around `turn = j;` with a red wavy line extending to the right.
- Red text `turn = i;` written next to the wavy line.

P_j

```
do {  
  while (turn == j)  
    CS  
    turn = i;  
} while (true);
```

Handwritten annotations:

- Red text `while (turn == j)` and `turn = i;` written over the printed code.
- Blue text `CS` written below the red `while` loop.

Solution to Critical-Section Problem

$n+1$ th

up to n

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a non-zero speed
 - No assumption concerning **relative speed** of the n processes

Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
 - Essentially free of race conditions in kernel mode

How to Implement Mutual Exclusion

XOR

- Three possibilities
- Sw • Application: programmer builds some method into the program
 - Hardware: special h/w instructions provided to implement ME
 - OS: provides some services that can be used by programmer
- All schemes rely on some code for
 - `enter_critical_section`, and
 - `exit_critical_section`

Application Mutual Exclusion

- Application Mutual Exclusion is
 - Implemented by the programmer
 - Hard to get correct, and very inefficient
- All rely on some form of busy waiting (process tests a condition, set a flag, and loops while the condition remains the same)

polling — busy waiting
interrupt

lock = 0 (unlocked)

Example

• Producer

entry (produce locked
if lock = 1 loop until lock = 0 unlocked
lock = 1
put in buffer
exit — lock = 0

— busy waiting

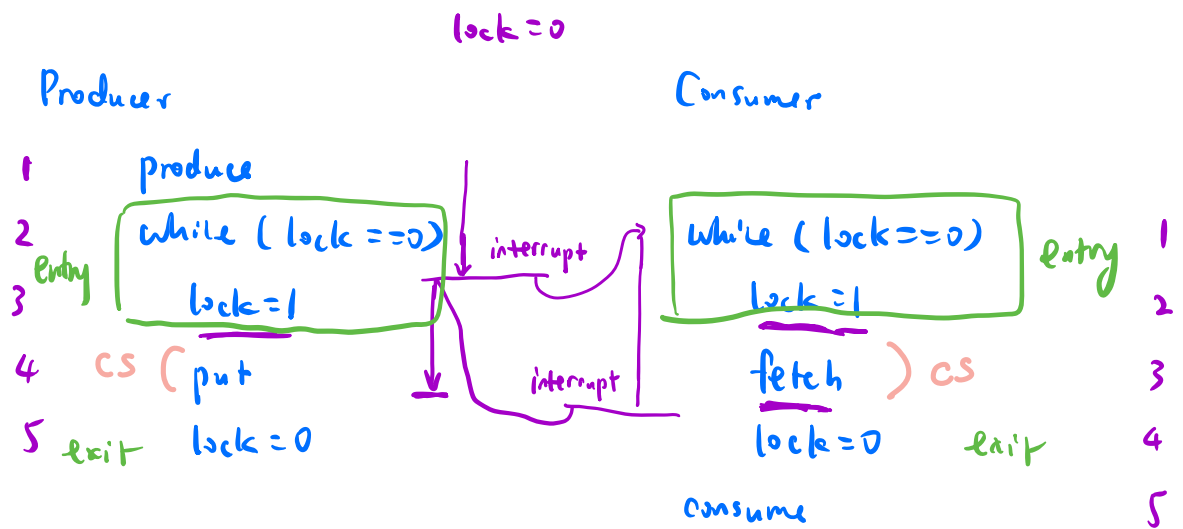
• Consumer

entry (if lock = 1 loop until lock = 0
lock = 1
get from buffer
exit — lock = 0
consume

— busy waiting

produce
while (lock == 0)
lock = 1
put) cs
lock = 0

while (lock == 0)
lock = 1
fetch) cs
lock = 0
consume



mutual exclusion violated

Locks


- A lock is a variable
- Two states
 - Available or free
 - Locked or held
- `lock()` : tries to acquire the lock
- `unlock()` : releases the lock which has been acquired by caller

non-preemptive

Disabling Interrupts

- One solution supported by hardware may be to use interrupt capability

```
do {  
    lock();  
    critical section  
    unlock();  
    remainder section  
} while (1);
```



```
void lock() {  
    DisableInterrupts();  
}  
void unlock() {  
    EnableInterrupts();  
}
```

⊕ Simple

⊖ interrupt lost

- Privileged instruction
- Multi-processor
- inefficient & slow

Disabling Interrupts (cont.)

- On a single CPU only one process is executed
- Concurrency is achieved by interleaving execution (usually done using interrupts)
- If you disable interrupts then you can be sure only one process will ever execute
- One process can lock a system or degrade performance greatly

Synchronization Hardware

- Many machines provide special hardware instructions to help achieve mutual exclusion
- The Test-And-Set (TAS) instruction tests and modifies the content of a memory word **atomically**
- TAS returns old value pointed to by **old_ptr** and updates said value to **new**

```
Int TestAndSet(int *old_ptr, int new) {  
    int old = *old_ptr;    //fetch old value at old_ptr  
    *old_ptr = new;        //store 'new' into old_ptr  
    return old;            // return the old value  
}
```

as a single unit

all or nothing

Mutual Exclusion with TAS

- Initially, lock's flag set to 0

```
typedef struct __lock_t {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    // 0 indicates that lock is available, 1 that it is held  
    lock->flag = 0;  
}  
  
void lock(lock_t *lock) {  
    while (TestAndSet(&lock->flag, 1) == 1)  
        ; // spin-wait (do nothing)  
}  
  
void unlock (lock_t *lock) {  
    lock->flag = 0;  
}
```

① unlocked (lock=0)

P_1 invokes lock()

↳ invokes TAS

↳ return 0
lock=1) *atomically*

② locked (lock=1)

P_2 invokes lock()

↳ invokes TAS

↳ return 1
lock=1) *atomically*
busy waiting

Busy Waiting and Spin Locks

- This approach is based on busy waiting
 - If the critical section is being used, waiting processes loop continuously at the entry point
- A binary “lock” variable that uses busy waiting is called a **spin lock**
 - Processes that find the lock unavailable “spin” at the entry
- It actually works (mutual exclusion)
- Disadvantages?
 - Fairness?
 - Performance?

Another HW Approach: Compare-And-Swap

- Test whether the value at the address specified by `ptr` is equal to `expected`
- If so, update the memory location pointed to by `ptr` with the new value. If not, do nothing.

```
int CompareAndSwap(int *ptr, int
expected, int new) {
    int actual = *ptr;
    if (actual == expected)
        *ptr = new;
    return actual;
}
```

```
void lock(lock_t *lock) {
    while (CompareAndSwap(&lock->flag,
0, 1) == 1)
        ; // spin
}
```

atomically

Hardware ME Characteristics

- Advantages

- Can be used by a single or multiple processes (with shared memory)
- Simple and therefore easy to verify
- Can support multiple critical sections

- Disadvantages

- Busy waiting is used
- Starvation is possible
- Deadlock is possible (especially with priorities)

Mutual Exclusion Through OS

- Semaphores
- Message passing IPC

Mutex Locks

mutual exclusion

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is **mutex** lock
- Protect a critical section by first `acquire()` a lock then `release()` the lock
 - Boolean variable indicating if lock is available or not
- Calls to `acquire()` ^{lock} and `release()` ^{unlock} must be **atomic**
 - Usually implemented via hardware atomic instructions
- But this solution requires **busy waiting**
 - This lock therefore called a **spinlock**

acquire() and release()

- ```
acquire() {
 while (!available)
 ; /* busy wait */
 available = false;;
}
```
- ```
release() {  
    available = true;  
}
```
- ```
do {
 acquire lock
 critical section
 release lock
 remainder section
} while (true);
```

# Semaphores

- Introduced by Edsger Dijkstra
- Motivation: Avoid busy waiting by blocking a process execution until some condition is satisfied
- Major advance incorporated into many modern operating systems (Unix, OS/2)
- A semaphore is
  - a non-negative integer
  - that has two valid operations

# Semaphore Operations

- Wait(s):

If  $s > 0$  then  $s := s - 1$

$S=0$  else block this process *sleep*

- Signal(s):

If there is a blocked process on this semaphore  
then wake it up

else  $s := s + 1$

- Wait() and Signal() originally called P() and V()


- prolaag: “probeer” in Dutch means “try”, and “verlaag” means “decrease” *wait*
- verhoog: Dutch for “increase” *signal*

# More on Semaphores

- Two types of semaphores
  - Binary semaphores can only be 0 or 1
    - Same as a mutex lock
  - Counting semaphores can be any non-negative integer
- Semaphores are an OS service implemented using one of the methods shown already
  - Usually by disabling interrupts for a very short time



# Producer – Consumer Problem: Solution by Semaphores



# Another Example

How many possible outputs?

3!

ABC  
ACB  
BAC  
BCA  
CAB  
CBA

- Three processes all share a resource on which
  - one draw an A
  - one draw a B
  - one draw a C
- Implement a form of synchronization so that the output appears ABC

*Process A*

```
think();

draw_A();
```

*Process B*

```
think();

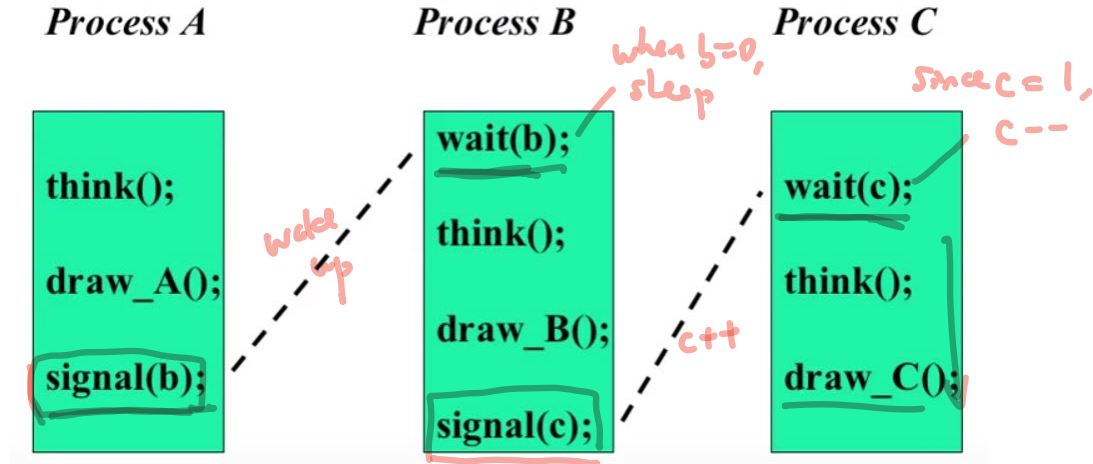
draw_B();
```

*Process C*

```
think();

draw_C();
```

- Semaphore  $b = 0$ ,  $c = 0$ ;



## Problem - Consumer

# Bounded-Buffer Problem

- We need 3 semaphores:

1. We need a semaphore mutex to have mutual exclusion on buffer access

ensure mutual exclusion

2. We need a semaphore full to synchronize producer and consumer on the number of consumable items

notify consumers how many items are there

3. We need a semaphore empty to synchronize producer and consumer on the number of empty spaces

notify producers how many empty slots are there

Counting semaphores

Avoid underflow ; Avoid overflow

# Bounded-Buffer - Semaphores

- Shared data

```
semaphore full, empty, mutex;
```

Initially:

```
full = 0, empty = n, mutex = 1
```

*no item to consume*      *all slots are available*      *— unlocked*

# Bounded-Buffer – Producer Process

wait(mutex);  
wait(empty);

```
do {
...
produce an item in nextp
```

```
...
wait(empty); if there is an empty slot to insert
wait(mutex); if there is any process accessing the buffer now
```

```
...
add nextp to buffer
```

```
...
signal(mutex);
signal(full);
} while (1);
```

X ↕ ↑

↑ ↕

# Bounded-Buffer – Consumer Process

do {

wait(full);

if there is any item to consume

wait(mutex);

if there is any process accessing the buffer now

...

remove item from buffer to **nextc**

...

signal(mutex);

signal(empty);

↕

has no effect

...

consume the item in **nextc**

...

} while (1);

X ↕

# Notes on Bounded-Buffer Solution

- Remarks (from consumer point of view)
  - Putting signal(empty) inside the CS of the consumer (instead of outside) has no effect since the producer must always wait for both semaphores before proceeding
  - The consumer must perform wait(full) before wait(mutex), otherwise deadlock occurs if consumer enters CS while the buffer is empty
- Conclusion: using semaphores is a difficult art



# Mutual Exclusion Problem: Starvation

- Definition

- Indefinitely delaying the scheduling of a process in favor of other processes

- Cause

- Usually a bias in a system scheduling policies

- Solution

- Implement some form of aging

# Another Problem: Deadlocks

- Two (or more) processes are blocked waiting for an event that will never occur
- Generally, A waits for B to do something and B is waiting for A
- Both are not doing anything so both events never occur

# Classical IPC Problems

## Producer-Consumer problem

- Bounded-buffer problem
- Readers and writers problem
  - Models access to a database (both read and write)
- Dining philosophers problem
  - Models processes competing for exclusive access to a limited number of resources such as I/O devices

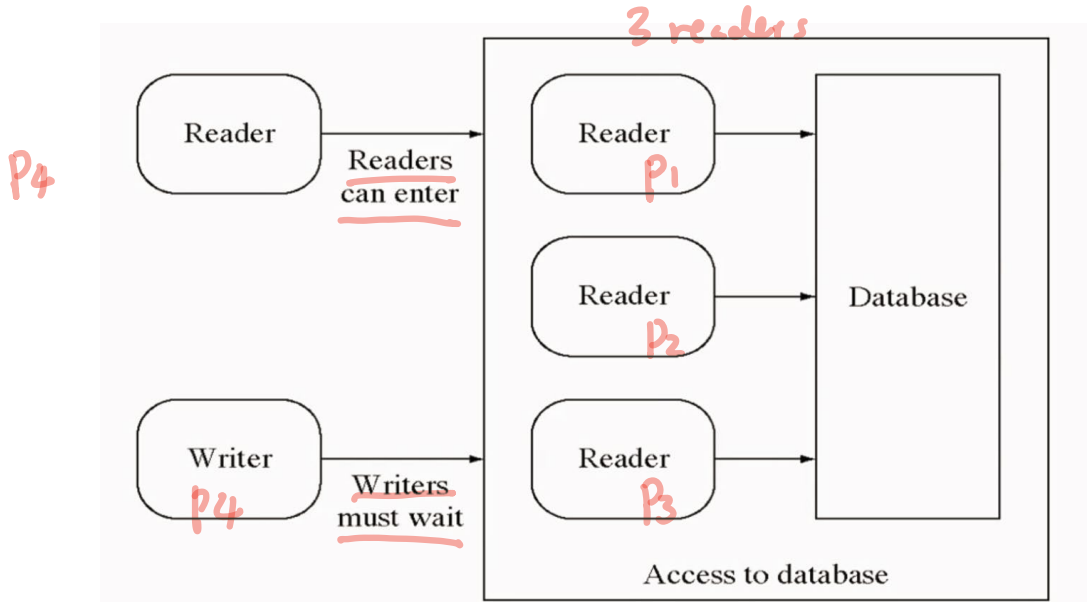
|          |                          |
|----------|--------------------------|
| producer | <u><math>r, w</math></u> |
| Consumer | <u><math>r, w</math></u> |
| <hr/>    |                          |
| reader   | $r$                      |
| Writer   | $w$                      |

# Readers-Writers Problem

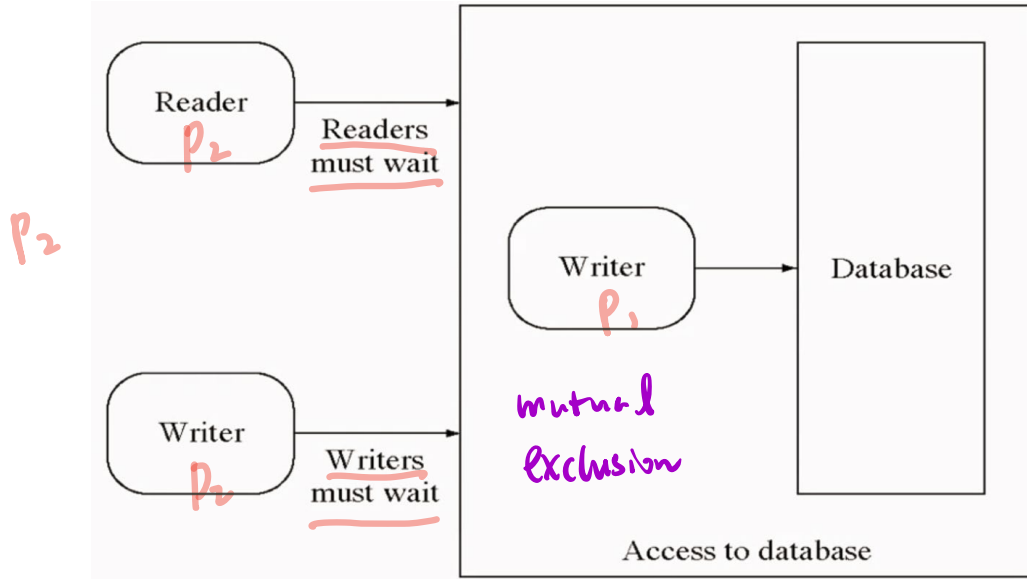
- Any number of reader activities and writer activities are running
- At any time, a reader activity may wish to read data
- At any time, a writer activity may want to modify the data
- Any number of readers may access the data simultaneously
- During the time a writer is writing, no other reader or writer may access the shared data

*mutual exclusion for writers*

# Readers-Writers with active readers



# Readers-Writers with an active writer

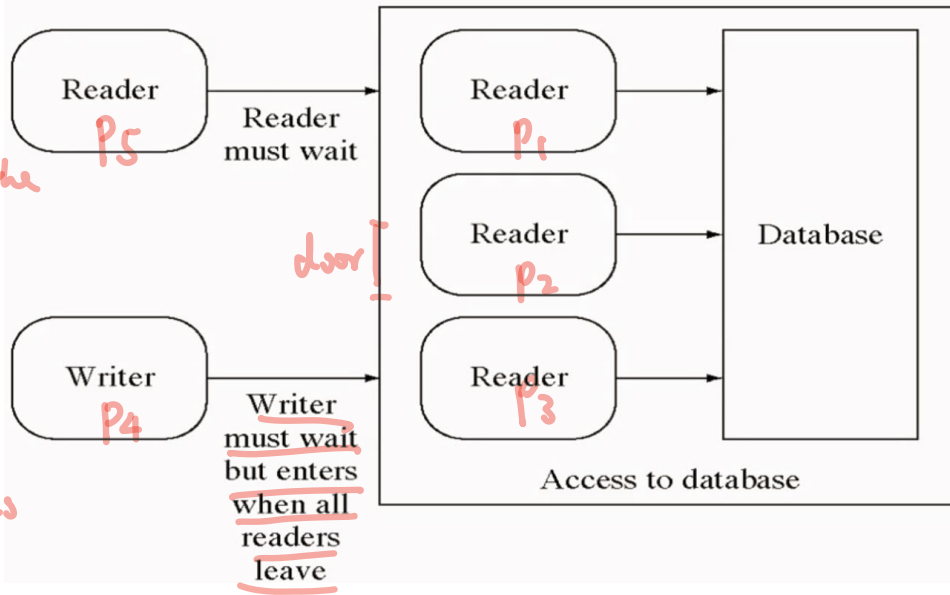


# Should readers wait for waiting writer?

3 active readers

Case 1:

P5 may skip the line, enter



Case 2:

P5 must wait until P4 leaves

# Readers-Writers problem

give priority to readers

1. The first readers-writers problem, requires that no reader will be kept waiting unless a writer has obtained access to the shared data
2. The second readers-writers problem, requires that once a writer is ready, no new readers may start reading *give priority to writers*
3. In a solution to the first case writers may starve; In a solution to the second case readers may starve.



# First Readers-Writers Solution

- **readcount** counter keeps track of how many processes are currently reading
- **mutex** semaphore provides mutual exclusion for updating readcount
- **wrt** semaphore provides mutual exclusion for the writers; it is also used by the first or last reader that enters or exits the CS

lock wrt when first reader enters

no writer  
can enter

unlock wrt when last reader leaves


writers  
can enter

solution to second RW problem  
requires five variables

no readers inside

# Dining Philosophers Problem

- Five philosophers are seated around a circular table
- In front of each one is a bowl of rice
- Between each pair of people there is a chopstick (fork), so there are five chopsticks
- It takes two chopsticks to eat rice, so while  $n$  is eating neither  $n+1$  or  $n-1$  can be eating



$$5 \text{ ppl} \times 2 \text{ chopstick} = 10$$

# Dining Philosophers Problem


- Each one thinks for a while, gets the chopsticks needed, eats, and puts the chopsticks down again, in an endless cycle
- Illustrates the difficulty of allocating resources among processes without deadlock and starvation
- The challenge is to grant requests for chopsticks while avoiding deadlock and starvation
- Deadlock can occur if everyone tries to get their chopsticks at once. Each gets a left chopstick, and is stuck, because each right chopstick is someone else's left chopstick

# Dining Philosophers Solution

- Each philosopher is a process
- One semaphore per fork
  - fork: array[0..4] of semaphores
- Initialization:
  - fork[i].count := 1
  - for i := 0..4

Process  $P_i$ :

```
repeat
 think;
 wait(fork[i]); — left
 wait(fork[i+1 mod 5]); — right
 eat;
 signal(fork[i+1 mod 5]); — right
 signal(fork[i]); — left
forever
```



2 ppl can eat  
at the same  
time.

4 ppl  
5 chopsticks

↓  
at least  
one person  
can get  
two sticks  
and eat

First attempt: deadlock if each philosopher starts by  
picking up his left fork (chopstick)!

# Dining Philosophers Solution

- Possible solutions to avoid deadlock:
  - Allow at most four philosophers to be sitting at the table at same time
  - Odd numbered philosopher picks up left fork first, even one picks up right fork



# Weakness of the Semaphore

- The user is expected to write wait and signal in the right order
  - The user must remember to execute signal for each exit
  - Calls may be spread throughout the program
  - The logic may demand that a process must check and signal his peers
- 
- There have been a wide number of alternatives proposed
  - Monitors are one common approach

# Summary

- We have seen two problems
  - Critical Sections – cannot both be modifying variable
  - Synchronization – must define ordering
- Often, our problems are a combination of the two
  - Readers/writers share storage, and readers should wait if there are writers waiting
- It is difficult to use semaphores correctly
- While there has been language support for Monitors for some time, standard UNIX still only supports semaphores (man sem\_open)

# Exit Slips

- Take 1-2 minutes to reflect on this lecture
- On a sheet of paper write:
  - One thing you learned in this lecture
  - One thing you didn't understand